

Homework 2 - Practical component

- This homework must be done and submitted to Gradescope in teams. You are welcome to discuss with students outside of your group but the solution submitted by a group must be its own. Note that we will use Gradescope's plagiarism detection feature. All suspected cases of plagiarism will be recorded and shared with university officials for further handling.
- The practical part should be coded in python (the only external libraries you can use are numpy and matplotlib) and all code will be submitted as a python file to Gradescope. To enable automated code grading you should work off of the template file given in this homework folder. Do not modify the name of the file or any of the function signatures of the template file or the code grading will not work for you. You may, of course, add new functions and any regular python imports.
- Any graphing, charts, derivations, or other practical report parts should be submitted to Gradescope included at the end of your report for the theoretical part of the homework.

You are of course encouraged to draw inspiration from what was done in lab sessions.

One-versus-all, L2 loss SVM

This part consists of the implementation of a one-versus-all, L2 SVM Loss, which is commonly used for multiclass classification. The L2 loss is differentiable and imposes a bigger penalty on points that violate the margin. In the one-versus-all (OVA) approach, we train m binary classifiers, one for

each class. At inference time, we select the class which classifies the test data with maximum margin.

Given a training set $S = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, where $\mathbf{x}_i \in \mathbb{R}^p$, $y_i \in \{1, \dots, m\}$ where p is the number of features and m is the number of classes, we would like to minimize the following objective function:

$$\frac{1}{n} \sum_{(\mathbf{x}_i, y_i) \in S} \sum_{j'=1}^m \mathcal{L}(\mathbf{w}^{j'}; (\mathbf{x}_i, y_i))^2 + \frac{C}{2} \sum_{j'=1}^m \|\mathbf{w}^{j'}\|_2^2$$

→ Reg

where

$$\mathcal{L}(\mathbf{w}^{j'}; (\mathbf{x}_i, y_i)) = \max\{0, 2 - (\langle \mathbf{w}^{j'}, \mathbf{x}_i \rangle) \mathbb{1}\{y_i = j'\}\}$$

and

$$\mathbb{1}\{y_i = j'\} = \begin{cases} 1 & \text{if } y_i = j' \\ -1 & \text{if } y_i \neq j' \end{cases}$$

In order to update the parameters \mathbf{w} of the SVM, gradient descent techniques are used. (Note: in the dataset provided with this assignment, the last element of each row is a dummy element with the value 1. That means there is no separate bias parameter b ; it is implicitly included in \mathbf{w} as the weight for the dummy element.)

The training set for this part can be downloaded from <https://drive.google.com/file/d/1Z5wfMe5DOLWTMDQXLGNbuw0Y7z6Kk31D/view?usp=sharing>.

The solution template file contains a function to load the data and do some preprocessing for you such as normalization. There are four files, one for the training features, training labels, test features and test labels.

Le fichier solution contient une fonction pour lire et effectuer quelques transformations sur les données telle que la normalisation. Quatre fichiers sont disponible, les *features* pour le jeu de données d'entraînement, un pour les cibles d'entraînement, un pour les *features* de test ainsi que les cibles.

<http://www.mit.edu/~9.520/spring09/Classes/class05.pdf> [5 pts] What is the derivative of the regularization term

$$\frac{C}{2} \sum_{j'=1}^m \|\mathbf{w}^{j'}\|_2^2$$

<https://maviccprp.github.io/a-support-vector-machine-in-just-a-few-lines-of-python-code/>

Refer CS229

with respect to w_k^j (the k th weight of the weight vector for the j th class)? Show all your work and write your answer in the report.

<http://www.cs.umd.edu/class/spring2017/cmsc422/slides0101/lecture12.pdf>

2. [10 pts] What is the derivative of the hinge loss term

$$\frac{1}{n} \sum_{(\mathbf{x}_i, y_i) \in S} \sum_{j'=1}^m \mathcal{L}(\mathbf{w}^{j'}; (\mathbf{x}_i, y_i))^2$$

with respect to w_k^j ?

Express your answer in terms of $\mathbf{x}_{i,k}$ (the k th entry of the i th training example \mathbf{x}_i).

Assume that

$$\frac{\partial}{\partial a} \max\{0, a\} = \begin{cases} 1 & \text{if } a > 0 \\ 0 & \text{if } a \leq 0 \end{cases}$$

(This is not exactly true: at $a = 0$, the derivative is undefined. However, for this problem, it's OK to make this assumption.)

3. [30 pts] Fill in the following in the code

- (a) [5 pts] `SVM.make_one_versus_all_labels`: Given an array of integer labels and the number of classes m , this function should create a 2-d array corresponding to the $\mathbb{1}\{y_i = j'\}$ term above. In this array, each row is filled with -1 , except for the entry corresponding to the correct label, which should have the entry 1. For example, if the array of labels is $[1, 0, 2]$ and $m = 4$, this function would return the following array: $[[-1, 1, -1, -1], [1, -1, -1, -1], [-1, -1, 1, -1]]$. The inputs are y (a numpy array of shape (number of labels,)) and m (an integer representing the number of classes), and the output should be a numpy array of shape (number of labels, m). For this homework, m will be 8, but you should write this function to work for any $m > 2$.
- (b) [5 pts] `SVM.compute_loss`: Given a minibatch of examples, this function should compute the loss function. The inputs are x (a numpy array of shape (minibatch size, 3073)), y (a numpy array of shape (minibatch size, 8)), and the output should be the computed loss, a single float.

- (c) [10 pts] `SVM.compute_gradient`: Given a minibatch of examples, this function should compute the gradient of the loss function with respect to the parameters \mathbf{w} . The inputs are X (a numpy array of shape (minibatch size, 3073)), y (a numpy array of shape (minibatch size, 8)), and the output should be the computed gradient, a numpy array of shape (3073, 8), the same shape as the parameter matrix \mathbf{w} . (Hint: use the expressions you derived above.)
 - (d) [5 pts] `SVM.infer`: Given a minibatch of examples, this function should infer the class for each example, i.e. which class has the highest score. The input is X (a numpy array of shape (minibatch size, 3073)), and the output is $y_inferred$ (a numpy array of shape (minibatch size, 8)). The output should be in the one-versus-all format, i.e. -1 for each class other than the inferred class, and $+1$ for the inferred class.
 - (e) [5 pts] `SVM.compute_accuracy`: Given an array of inferred labels and an array of true labels, this function should output the accuracy as a float between 0 and 1. The inputs are $y_inferred$ (a numpy array of shape (minibatch size, 8)) and y (a numpy array of shape (minibatch size, 8)), and the output is a single float.
4. [5 pts] The method `SVM.fit` uses the code you wrote above to train the SVM. After each epoch (one pass through the training set), `SVM.fit` computes the training loss, the training accuracy, the test loss, and the test accuracy.

Plot the value of these four quantities for every epoch for $C = 0.1, 1, 30, 50$. Use 200 epochs, a learning rate of 0.001, and a minibatch size of 5000.

You should have four plots: one for each quantity, with the curves for all four values of C . Include these four plots in your report.