



- Published Aug 10
- 👺 Fork of Unsupervised Learning by 🦣 Kris Sankaran

## **Unsupervised Learning**

IFT6758, Fall 2020

Reading: ISLR sections 10.1, 10.2 and PDS pg. 462 - 476

#### What is unsupervised learning?

- Supervised learning: Learn some mapping  $x_i o y_i$
- Unsupervised learning
  - $\circ$  Usual definition: "Exploring the data  $x_i$ "
- Less orthodox interpretations
  - Learning with hidden labels (clusters: missing classes)
  - Data compression for human consumption
  - o Generation of derived data for human consumption

#### **Common Themes**

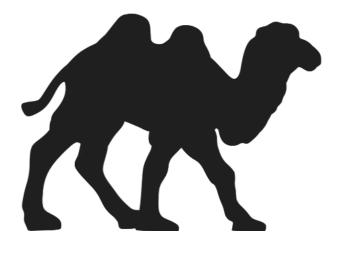
. Much handauta analusta

- Much narger to evaluate
- Most methods can be categorized as either,
  - Dimensionality reduction: Few features that summarize many
  - Clustering: Few "prototypes" that are representative of whole dataset
- We'll review canonical examples,
  - PCA (dimensionality reduction)
  - K-means and Hierarchical Clustering (clustering)

## **PCA**What is this?



**PCA**What is this?



Credit for the idea: Prof. Julie Josse

#### **PCA**

Idea: Certain views of high-dimensional data are more informative than others.

Can you find a low-dimensional representation with as much variation as possible?

Can you find a low-dimensional representation with as much variation as possible?

To implement the idea,

- What will be the candidate family of low-dimensional representations?
- How will we choose one of the many candidates?

## **Candidates: Linear Mixings**

Can you find a **low-dimensional representation** with as much variation as possible?

• For a representation, consider linear combinations of high-dimensional vectors,

$$egin{aligned} z_i &= \sum_{j=1}^p arphi_{1j} x_{ij} \ &= arphi_1^T x_i \end{aligned}$$

•  $\varphi_1$  is a free parameter. E.g., if  $\varphi_1 = \left(\frac{1}{p}, \dots, \frac{1}{p}\right)$ , then we summarize  $x_i$  by averaging over its coordinates

#### **Selection Criteria: Maximal Variance**

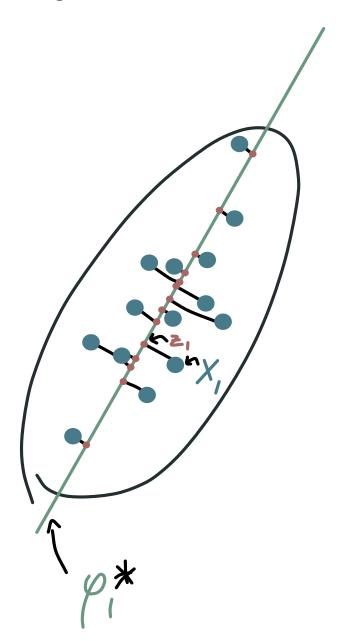
Can you find a low-dimensional representation with **as much variation** as possible?

• The  $z_i$ 's should be as spread out as possible:

$$ext{maximize}_{arphi_1} rac{1}{n} \sum_{i=1}^n z_i^2$$

 $\circ$  Subject to constraint  $\|arphi_1\|^2=1$ .

See the example here. The red arrow is  $\varphi_1$  and the points along the 1D axes are the associated  $z_i$ 's.





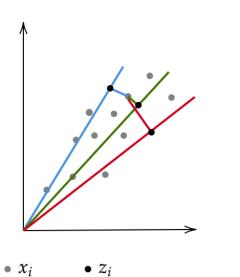
Can you find a low-dimensional representation with **as much variation** as possible?

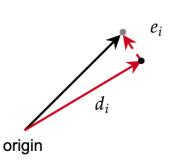
• The  $z_i$ 's should be as spread out as possible:

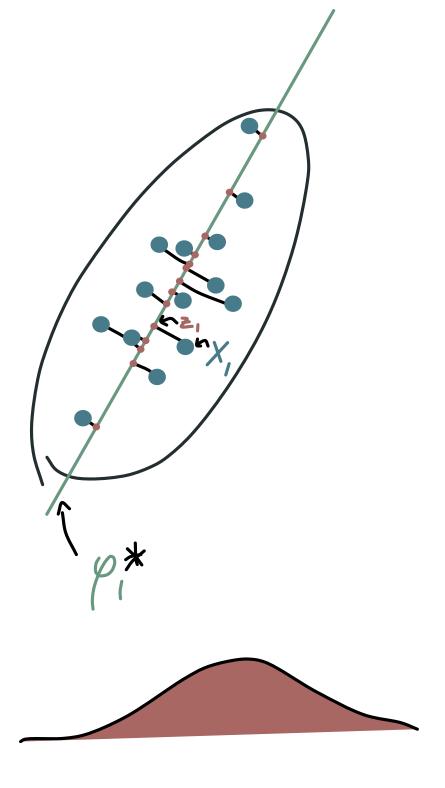
$$ext{maximize}_{arphi_1}rac{1}{n}\sum_{i=1}^n z_i^2$$

 $\circ$  Subject to constraint  $\|arphi_1\|^2=1$ .

See the example here. The red arrow is  $\varphi_1$  and the points along the 1D axes are the associated  $z_i$ 's.



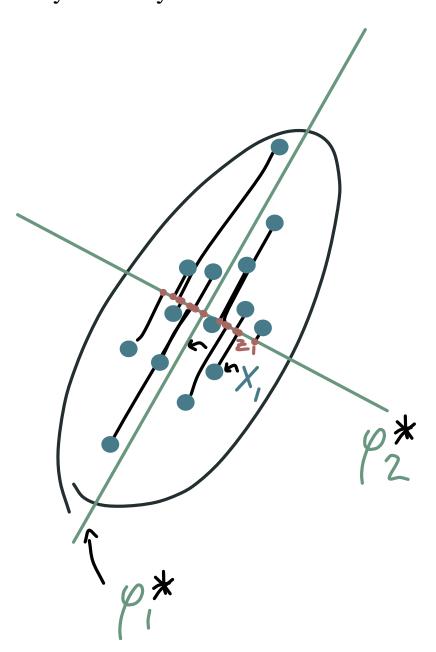




## Second, third, ... PCA directions

- Once you find  $\varphi_1$ , you can find a "second" direction  $\varphi_2$
- Found by solving the exact same optimization, but with a

- new constraint that it's at 90 degrees to the previous directions
- Interpretation: PCA is finding a new, better coordinate system for your data



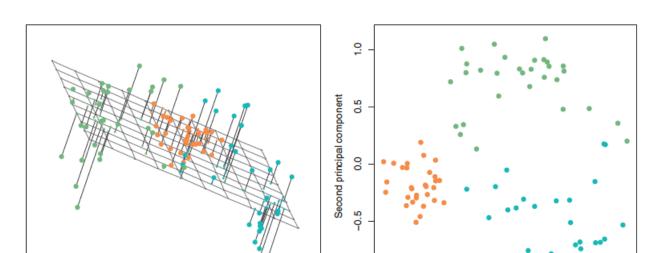
**Semantics** 

## **Alternative Interpretation: Linear Approximation**

- The first K directions in PCA find the best Kdimensional linear approximation (using sum of squared
  error to measure approximation quality).
- This means it's fair to say

$$x_{ij}pprox \sum_{k=1}^K z_{ik} arphi_{kj}$$

- Or, in matrix notation,  $x_i pprox \Phi z_i$ , where  $\Phi$  concatenates the  $\varphi_k$ 's vertically
  - $\circ \ x_i$  is a mixture of the components  $arphi_k$  with weights  $z_{ik}$  .



•



**Example with two principal components** 

We'll practice reading PCA plots using USArrests --- a little dark but you will see more of such *breaking bad* themed examples in ISLR

- $z_i$ 's give the states' coordinates
- Can interpret components by looking at how variables contribute. Variable j is plotted at  $(\varphi_{1j}, \varphi_{2j})$ .
  - E.g., the second PC mostly captures variation related to urban population

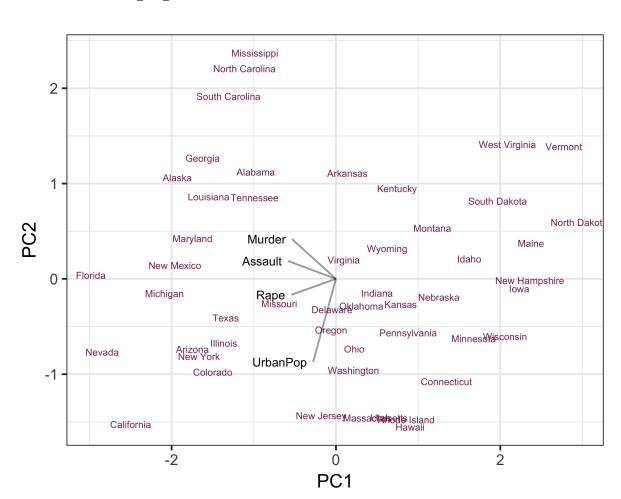
-0.5

-1.0

0.5

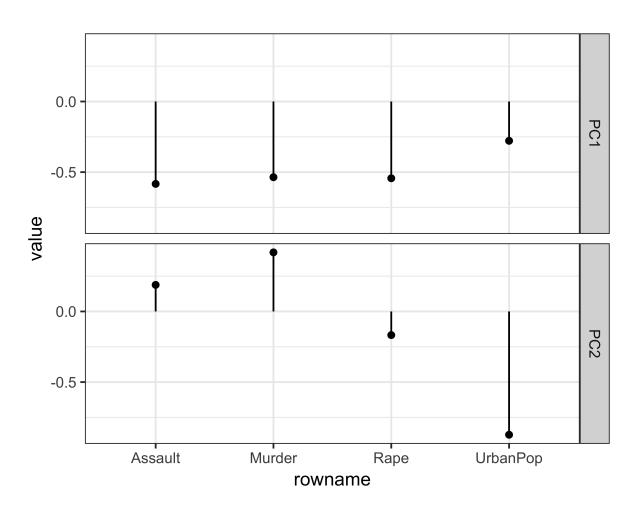
First principal component

1.0



## **Biplots**

Arrows come from the  $\varphi_1$  and  $\varphi_2$ . The (x, y)-coordinate of the arrows comes from viewing these PCs in 2D.

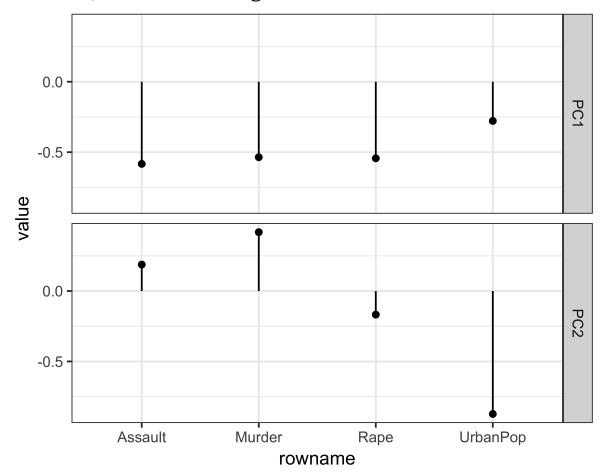


## **Biplots**

• The coordinate of  $r_i$  on the hiplot is  $(x_i, x_i)$ 

The coordinate of  $\omega_i$  on the diplot is  $(z_{i1}, z_{i2})$ 

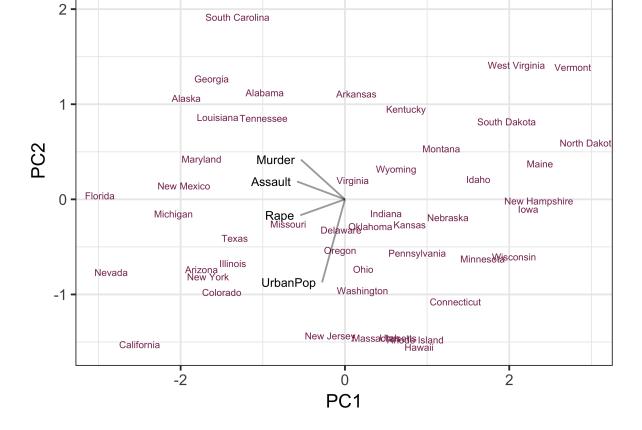
• Since  $x_i \approx z_{i1}\varphi_1 + z_{i2}\varphi_2$ , they have large values for variables with large loadings in the coordinate directions where  $x_i$  is farther along



## **Biplots**

- For example, California  $pprox -2.5 arphi_1 1.52 arphi_2$
- Since  $\varphi_1$  puts negative weight on the crimes, California has more than the average # of crimes  $(-\times -=+)$
- Since  $\varphi_2$  puts negative weight on urban population, California has larger than the average population





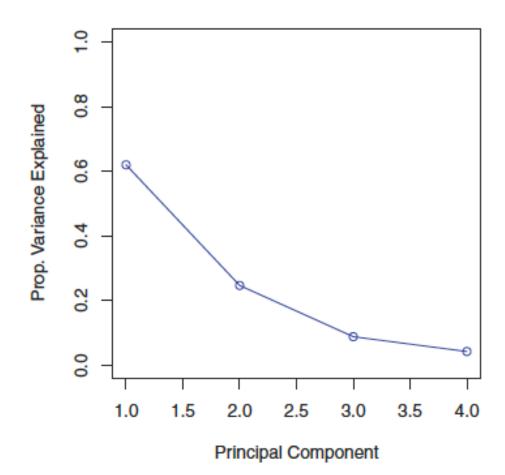
#### **Explained Variation**

• Amount explained by  $k^{th}$  component,

$$rac{\sum_{i=1}^{n}\|z_{ki}\|^2}{\sum_{i=1}^{n}\|x_i\|^2}$$

- If no directions are preferred, get  $pprox rac{1}{p}\%$  everywhere
- How to compute?
  - Consider the covariance matrix of the data. Check total variance. Now look at the transformed covariance matrix and find the ratio of variances of each component.

- Exercise: What would the plot below look like if the data were shaped like...
  - a pancake (two long directions, one short one)
  - a cigar (one long direction, two short ones).



#### Things to watch out for

- Even though the method is easy to run, there are lots of potential issues,
  - Variables might be at different scales, and there might be ambiguity about whether to rescale them
  - o The directions are only unique up to sign

• Choosing *K* is tricky (although might not be crucial)

#### Clustering

Idea: Partition the observations, so that those that are similar to each other appear together

Look for homogeneous subgroups in your heterogeneous data.

#### **Formalization**

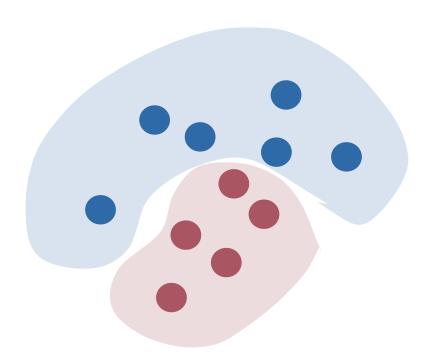
A partition  $C_1, \ldots, C_k$  is a collection of subsets satisfying,

• Each sample is in a subset:

$$\cup_{k=1}^K C_k = \{1,\ldots,n\}$$

• Subsets are disjoint: For any pair,

$$\alpha \circ \alpha = \alpha$$



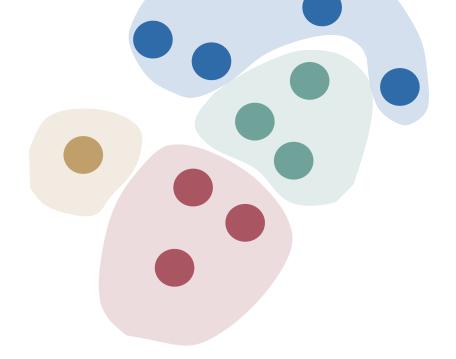
A partition  $C_1, \ldots, C_k$  is a collection of subsets satisfying,

• Each sample is in a subset:

$$\cup_{k=1}^K C_k = \{1,\ldots,n\}$$

• Subsets are disjoint: For any pair,

$$C_k \cap C_{k'} = \emptyset$$



A partition  $C_1, \ldots, C_k$  is a collection of subsets satisfying,

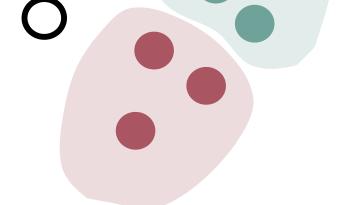
• Each sample is in a subset:

$$\cup_{k=1}^K C_k = \{1,\ldots,n\}$$

• Subsets are disjoint: For any pair,

$$C_k \cap C_{k'} = \emptyset$$





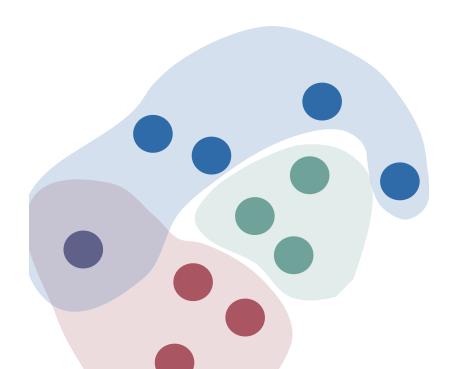
A partition  $C_1, \ldots, C_k$  is a collection of subsets satisfying,

• Each sample is in a subset:

$$\cup_{k=1}^K C_k = \{1,\ldots,n\}$$

• Subsets are disjoint: For any pair,

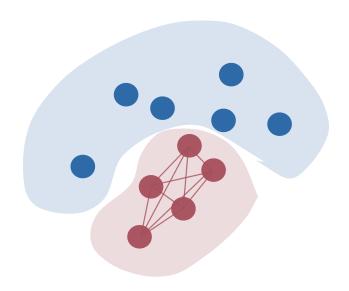
$$C_k\cap C_{k'}=\emptyset$$



We don't want just any partition, but the one that minimizes within group variation, which we'll call W.

$$\mathop {\operatorname {minimize}}\limits_{{C_k}} \sum\limits_{k = 1}^K {W\left( {{C_k}} 
ight)}$$

Usually, we use  $W\left(C_k
ight) = rac{1}{|C_k|} \sum_{i,i' \in C_k} \|x_i - x_i'\|^2$  .



#### **Algorithm**

This is a combinatorial optimization problem, and finding the global optimum is computationally challenging.

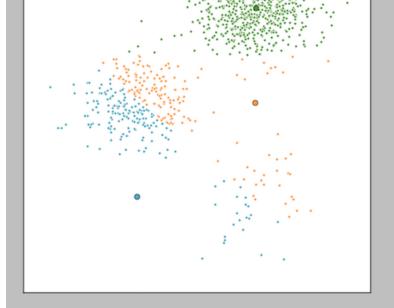
However the following algorithm usually finds good local optima,

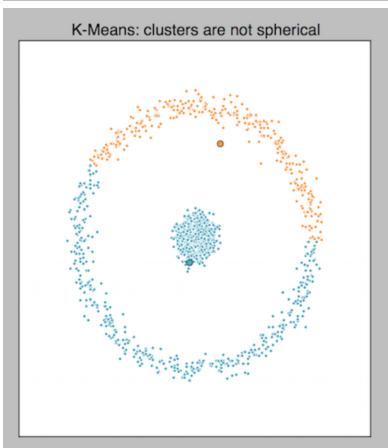
- 1. Arbitrarily assign each  $x_i$  to one of the clusters,  $C_1, \ldots, C_K$ .
- 2. Iterate until convergence,
  - a. Compute the mean  $\bar{x}_k$  of the points in  $C_k$ .
  - b. Reassign the points  $x_i$ , so they are put in the cluster whose centroid they are closest to.

Here is a nice demo. The procedure reduces the criterion under study at each step, which means we will converge to a local optimum.

#### Limitations of K-means

K-means often does not do a good job in clustering when there are variations in density, nonspherical shapes of clusetrs and outliers.

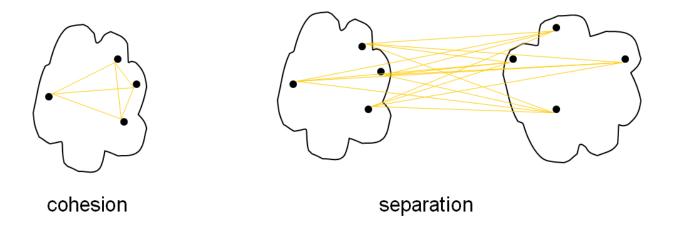




#### How to evaluate K-means results?

- Supervised evaluation: use a pre-classified dataset as a benchmark
  - Given the knowledge of the ground truth class assignments of the samples, use some metric, e.g., homogeneity, completeness, v-score, to evaluate the goodness of the the clustering (higher score is better)

- Python: from **sklearn** import **metrics**; metrics.homogeneity\_score(labels\_true, labels\_pred)
- Unsupervised evaluation: minimize intra-cluser distance (maximize cohesiveness) and maximize inter-clusetr distance (separation)



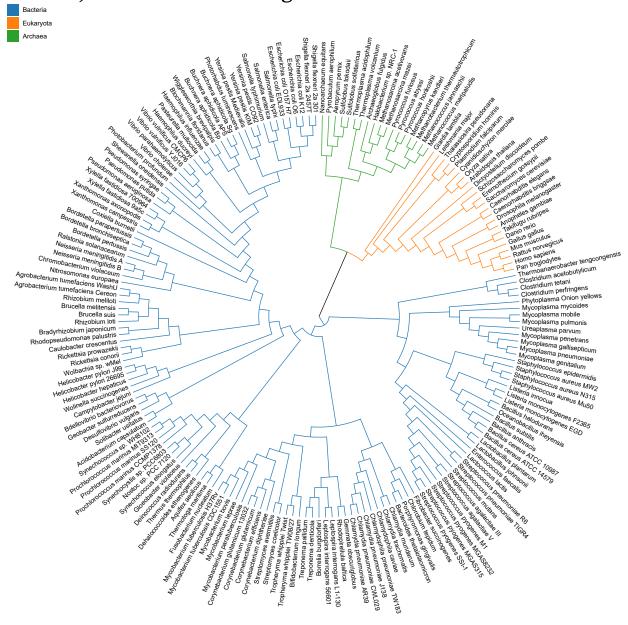
Details: Cluster Analysis Using K-means Explained

## **Hierarchical Clustering**

- *K* controls the "magnification" at which we do the clustering.
- What if we could do the clustering at many different scales, all at once?

## **Hierarchical Clustering**

Behold, the cluster dendrogram.

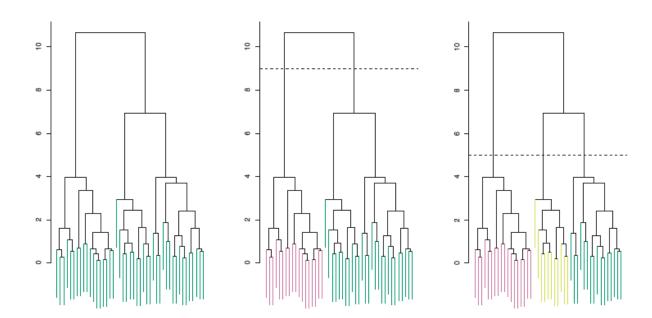


## **Interpretation**

• Samples which are similar to each other are put on the

same subtree.

- Pairs of samples that are very similar to one another share very recent common ancestors
  - Beware: Samples can be close by at the leaves without being close in the subtree sense
- You can get a standard clustering by "cutting" tree at some horizontal level

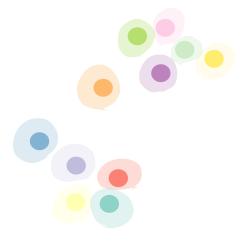


#### That's cool, how do I make it?

- These trees are informative. We'd like an automated procedure for creating them.
  - Agglomerative: *bottom-up*
  - Divisive: *top-down*

#### **Agglomerative**

a. Initialize: Associate each point with a cluster  $C_i := \{x_i\}$ b. Iterate until only one cluster: Look at all pairs of clusters. Merge the pair  $C_k$ ,  $C_{k'}$  which are the most similar.



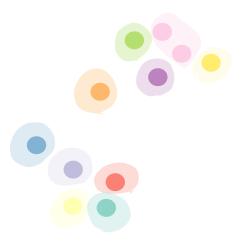
# 

#### That's cool, how do I make it?

• These trees are informative. We'd like an automated procedure for creating them.

#### **Agglomerative**

a. Initialize: Associate each point with a cluster  $C_i := \{x_i\}$  b. Iterate until only one cluster: Look at all pairs of clusters. Merge the pair  $C_k$ ,  $C_{k'}$  which are the most similar.





#### That's cool, how do I make it?

• These trees are informative. We'd like an automated procedure for creating them.

#### **Agglomerative**

a. Initialize: Associate each point with a cluster  $C_i := \{x_i\}$  b. Iterate until only one cluster: Look at all pairs of clusters. Merge the pair  $C_k$ ,  $C_{k'}$  which are the most similar.



#### That's cool, how do I make it?

• These trees are informative. We'd like an automated procedure for creating them.

#### **Agglomerative**

a. Initialize: Associate each point with a cluster  $C_i := \{x_i\}$ 

Merge the pair  $C_k$ ,  $C_{k'}$  which are the most similar.



#### That's cool, how do I make it?

• These trees are informative. We'd like an automated procedure for creating them.

#### **Agglomerative**

a. Initialize: Associate each point with a cluster  $C_i := \{x_i\}$  b. Iterate until only one cluster: Look at all pairs of clusters. Merge the pair  $C_k$ ,  $C_{k'}$  which are the most similar.

